

Rep Smarter, Not Harder: AI Hypertrophy Coaching with Wearable Sensors and Edge Neural Networks

1st Musa Azeem

College of Engineering and Computing
University of South Carolina
Columbia, SC, USA
mmazeem@email.sc.edu

2nd Grant King

College of Engineering and Computing
University of South Carolina
Columbia, SC, USA
gking@email.sc.edu

3rd Savannah Noblitt

College of Engineering and Computing
University of South Carolina
Columbia, SC, USA
snoblitt@email.sc.edu

Abstract—abstract

I. INTRODUCTION

brief introduction...

A. Background

Skeletal muscle hypertrophy is a well-recognized physiological response to resistance training (RT), primarily driven by mechanical tension applied to muscle fibers [1]. To promote meaningful hypertrophic gains, individuals are generally advised to perform resistance training sets to near momentary muscular failure: the point at which no more repetitions can be performed with correct form. This idea is often described by the number of repetitions in reserve (RiR) left upon completing a set, which is the estimated number of additional repetitions an individual could have completed before reaching momentary muscular failure [1]. Adjusting RT based on RiR can be beneficial across various training populations. Bodybuilders training for hypertrophy, for example, may choose to leave a few repetitions in reserve to reduce accumulated fatigue while still maintaining progress. Likewise, athletes training for strength and power can strategically manage RiR to optimize performance outcomes [1], [2]. Lastly, incorporating RT into physical therapy has demonstrated positive effects on muscle function and recovery in rehabilitation settings [3]. In such cases, adjusting repetitions in reserve can play a key role in managing the balance between promoting progress and avoiding reinjury.

B. Related Work

To the best of our knowledge, there is no existing study that focuses on directly predicting RiR in real time using wearable sensors, which is an area this work aims to advance. However, several studies have explored the use of wearable sensor data for related predictive tasks.

Albert et al. [4] investigated the use of machine learning to predict Ratings of Perceived Exertion (RPE) based on physiological and motion data. RPE is a subjective 1–10 scale reflecting how hard an individual feels they exerted themselves during a set. Sixteen male participants with at least one year of strength training experience performed squats on a flywheel machine, which provides resistance by requiring users to

decelerate a rotating flywheel during the eccentric phase. Each participant completed four squat series, with each series consisting of three 12-repetition sets, interspersed with short rest periods. Sessions lasted one hour, with ECG recorded for 20 minutes post-exercise and RPE scores collected 15 minutes after the final set. Data were collected via six IMUs and one ECG sensor, all streamed to an Android app via Bluetooth. IMU data were filtered using a Butterworth filter, and heart rate variability (HRV) metrics were extracted from ECG data. A total of 288 features were generated (6 sensors $\times$ 6 axes $\times$ 8 statistical features). Windowed data from IMUs were synchronized with ECG-derived HRV features, and feature selection was applied before model training. Four regression models were evaluated: Support Vector Regression (SVR) with both RBF and linear kernels, Random Forest Regression (RFR), and Gradient Boosted Regression Trees (GBRT). Due to class imbalance in RPE scores, resampling was performed. Performance was assessed using MSE, RMSE, MAPE, and Pearson correlation (r). Out of 2162 repetitions across 181 sets (with 2 subjects dropping out due to fatigue), the best performance came from the GBRT model using 5-second windows with 50% overlap, achieving a Pearson $r = 0.89$. This study addresses a closely related problem by using wearable sensors to estimate rate of perceived exertion (RPE), which is conceptually linked to repetitions in reserve (RiR) as its approximate inverse. Although it relies on ECG data in addition to IMUs, its use of windowing strategies and feature extraction techniques reflects common practices in sensor-based effort estimation.

Jiang et al. [5] modeled fatigue onset on a per-repetition basis across three exercises—squat, high knee jack, and corkscrew toe-touches. Fourteen healthy adults (12 men, 2 women) performed sets to exhaustion, with a minimum of three sets per exercise. Each set contained five repetitions, and RPE was recorded after each set. The number of sets varied by participant. Multimodal sensor data was collected using 17 IMUs, 32 reflective motion capture markers, and a force plate. Force plate features included center of pressure (used to identify reps) and vertical ground reaction forces. IMU accelerations and angular velocities were also extracted. Sensor data were synchronized using dynamic time warping, segmented into individual reps, and normalized via cubic spline interpo-

lation to fixed-length windows. Noise was reduced using Butterworth filtering, and class imbalance was addressed through upsampling. Two models were trained: Random Forest Regression and a CNN with three 2D convolutional layers (kernel size [10, 3], 32–64–128 filters, ReLU activation, 0.8 dropout) followed by a fully connected layer with 128 units and sigmoid activation. Models were trained in a subject-dependent manner. The CNN outperformed Random Forest on squats and high knee jacks, while performance was similar for corkscrew toe-touches. Peak Pearson correlation coefficients reached 89% (squat), 93% (jack), and 94% (corkscrew), with performance improving as participants completed more sets. This study focuses on rep-by-rep fatigue estimation using IMU-based deep learning with a convolutional neural network (CNN) architecture. Its use of subject-dependent models reflects a common approach in sensor-based fatigue research, though it also highlights the challenge of achieving generalizability across individuals.

Elshafei and Shihab [6] focused on detecting bicep muscle fatigue during concentration curl exercises using wearable sensors. A key distinction here is the prediction of fatigue, a decrease in the mechanical capabilities of a muscle, rather than muscular failure. They collected data from 20 middle-aged volunteers (ages 29–33, BMI 18–28), all with at least one year of gym experience and no history of chronic diseases or surgeries. Each participant performed concentration bicep curls while wearing a 9-axis Inertial Measurement Unit (IMU) on the wrist and an Apple Watch Series 4 to monitor heart rate. Each repetition was labelled as fatigued or not based on reported RPE and heartrate. The study faced challenges such as selecting an appropriate dumbbell weight — 4.5 kg was found optimal for balancing session length and fatigue induction — and addressing the subjectivity of RPE by averaging self-reported RPE with heart rate-derived estimates in cases of discrepancy between the two. They observed that fatigue reduced the biceps’ angular velocity, increasing completion time in later sets. From the dataset, 33 features were extracted and reduced to 16 significant ones. Five machine learning models were evaluated: Generalized Linear Models (GLM), Logistic Regression (LR), Random Forests (RF), Decision Trees (DT), and Feedforward Neural Networks (FNN). The two-layer FNN achieved the highest accuracy, with 98% for subject-specific models and 88% for cross-subject models. This study demonstrates the feasibility of using wearable IMU sensors to detect muscle fatigue, a task highly related to our own. It also highlighted the importance of feature selection and the potential of neural networks in modeling fatigue across different individuals.

Velloso et al. [7] explored qualitative activity recognition by assessing the execution quality of weightlifting exercises using wearable sensors. They aimed to detect mistakes in exercise form and provide real-time feedback to users. The study involved participants performing weightlifting exercises while equipped with an array of IMU sensors to capture motion data. Two approaches were evaluated for qualitative activity recognition: classification of wearable sensor data and

a model-based method, where performances were compared to an ideal-form model. For the classification approach, which we find relevant to our study, six young male participants performed 10 dumbbell curl repetitions in five different ways: 4 with varying mistakes in the exercise form and 1 with correct form. These repetitions were recorded by four 9-axis IMU sensors on the wrist, bicep, lifting belt, and dumbbell. An ensemble of Random Forest models was trained to predict the class a repetition belonged to. The model achieved an accuracy of 98.03% using 10-fold cross validation across subjects and an accuracy 78.2% using leave-one-subject-out cross validation. Through this study, a window size of 2.5s with an overlap of 0.5s was determined to be the best for this classification task. The IMU features with the highest correlation to the classes were also identified.

Dang et al. [8] investigated muscle fatigue classification during force-relaxation cycles using surface electromyography (sEMG) signals collected from the biceps and triceps of 10 healthy young men who had not engaged in strenuous activity for 3 days. Their goal was to classify muscle activity into one of three fatigue states: non-fatigued, semi-fatigued, or fatigued. They designed a dual-lead sEMG sensing device which consisted of a main template, lead wires, adjustable gain, and batteries. This non-invasive device was attached to the biceps and triceps to record sEMG signals during force-relaxation cycles. The level of force exerted during these cycles varied purposefully to induce the 3 different fatigue states. A total of 240 sets of sEMG sequence data were collected — 80 for each of the three fatigue states. This collected dataset was then utilized to evaluate four different machine learning algorithms (DTW-KNN, FNN, LSTM, and Attention-LSTM). Of these, the Attention-LSTM performed the best, it achieved a classification accuracy of 93.5% and a classification assessment time of 3.73 seconds. This study highlights an advancement in muscle fatigue classification by demonstrating that inputs such as sEMG signals can effectively complement or substitute IMU data in the implemented machine learning models.

While prior research has addressed fatigue classification, perceived exertion estimation, and lifting quality, to our knowledge, no existing study has proposed a complete, sensor-driven pipeline for real-time prediction of repetitions in reserve. This remains an open area of research with substantial opportunity for further exploration and advancement.

C. Contributions

In this work we explore the novel task of predicting RiR in real time using wearable, easy-to-apply sensors. As a first step toward this goal, we present a novel neural network-based pipeline for detecting “near-failure” repetitions in real time using a wrist-mounted inertial measurement unit (IMU). Our contributions are as follows:

- We introduce a new dataset of wrist IMU data collected during preacher curls performed to momentary muscular failure by 13 diverse participants. The dataset includes

over 600 annotated repetitions with varied form, weight, and number of reps.

- We develop a segmentation model to detect the end of each repetition in real time.
- We develop a classification model to identify near-failure repetitions, defined as repetitions performed with two or fewer remaining repetitions in reserve ($RiR \leq 2$).
- We demonstrate that our models can be deployed on edge hardware using a Raspberry Pi 5, validating their potential for real-time, low-power operation in various gym settings.

II. MATERIALS AND METHODS

A. Data Collection

1) *Participants*: Participants for data collection consisted of this work's authors and volunteers at three different gyms. There were 13 participants in total (9 male, 4 female). All participants were about 18-25 years old and varied in weight, height, strength, and resistance training experience.

2) *Equipment*: Participant preacher curls were recorded using a custom-made recording device, consisting of a Raspberry Pi 5 and the Adafruit MPU-6050 6-axis inertial measurement unit (IMU), capturing both accelerometer and gyroscope data. Although only IMU data was used here, the device, shown in Figure 1 was designed with the intention and capability of recording data from multiple sensors (such as heartbeat sensor and sEMG). The Raspberry Pi read data from the IMU at 100 Hz, and was connected to and accessed by a laptop. Three different preacher curl machines at the three locations were used for data collection: (TODO).




Fig. 1. **Data collection device.** The Raspberry Pi (bottom) reads, processes, and stores data from the IMU (top right). The sEMG sensor (bottom right) was not used due to hardware issues. During data collection, the Pi was powered by a portable battery pack. A box was used to protect the Raspberry Pi from damage, and the IMU was taped to participants' wrists.

3) *Procedure*: During data collection, the IMU was taped to each participant's left wrist. The axes of the sensor were oriented as shown in Figure 2. Each participant was instructed to perform sets of preacher curls with consistent form, taking each set to momentary muscular failure. Participants were free to select the weight they lifted and decide how many sets they performed. Consequently, collected data varied in the

amount of weight lifted, the number of repetitions per set, and how many sets were completed in a session. Participants also introduced natural variation in form and range of motion (e.g. locking out their arms at the bottom of the curl, duration of pause at the bottom, height of the concentric phase). Only sets taken to failure with proper form were admitted to this work's dataset.




Fig. 2. **IMU orientation.** The IMU was taped to the participant's wrist with the axes oriented as shown. Image source [9].

During data collection, annotations were recorded for future labelling. This was facilitated by a custom Python server running on the Raspberry Pi. When prompted to start recording, the server read and stored data from the IMU in CSV format at approximately 100 Hz. The server broadcasts an interactive data collection web application for access via laptop over Wifi or Ethernet. This data collection dashboard was used to enter participant metadata and location metadata, start the recording of IMU data, manually record timestamp markers to annotate the end of each repetition, and stop the recording. The timestamp markers were later used to generate data labels.

In total, 13 participants performed 68 sets of preacher curls, with an length of 36 seconds and 9.3 curl repetitions. This resulted in 40 minutes of data consisting of 631 total reps.

B. Data Processing

Several steps were taken to process the raw IMU data and prepare it for model training.

1) *Data Cleaning*: First, to resolve fluctuations in data sampling rate caused by the Raspberry Pi's CPU load, the 6-channel timeseries data was first interpolated to exactly 100 Hz. Next, the data was smoothed with a 15-data point (150 ms) moving average filter to remove signal noise. The data annotations were adjusted as well. Manual inspection of the timeseries signals was used to adjust any incorrect end-of-rep annotations to the true end of the rep. A full set of curls after cleaning, along with the data annotations, can be seen in Figure 3.

2) *Splitting & Windowing*: To facilitate model training, the collected data was split into train and validation splits. About 80% of the recorded preacher curl sets (53) were added to the train set. The remaining 20% of sets (15) were added to the validation split. Next, the data in each split was windowed into 256 data-point (2.56s) segments, with a stride of 2 data




Fig. 3. **Full set of preacher curls.** The accelerometer (top) and gyroscope (bottom) data from a single set of preacher curls, with the end-of-rep annotations shown as red vertical lines. The set lasted 26.2 seconds and consisted of 8 repetitions to failure. Red numbers indicate the Repetitions in Reserve at each point during the set. The last rep, at which point the participant could no longer lift the weight, is a rep taken to failure, with an RiR of 0.

points (20 ms). Each window was labelled based on the end-rep-marker annotations, as will be explored. In total, about 90k training windows and 26k validation windows were generated.

Splitting data at the set level guarantees that no overlap exists between the data in the training and validation set, ensuring that the model is not trained on data it will be evaluated on. Windows of 2.56 seconds were chosen to balance the receptive field of the model with real time capabilities. Windows of 2.56 ensure models had the context needed to capture an entire rep, which are usually less than 2.56 seconds. At the same time, the 2.56 second windows are small enough to allow for real time inference on edge devices. A small stride was chosen to maximize the number of training samples available to the model.

3) *Data Augmentation*: To simulate additional data, the training data was augmented with several techniques. First, training samples were randomly stretched and cropped in the time dimension by a factor of 1 to 1.5, simulating faster reps. Next, the amplitude of the training windows were randomly scaled by a factor of 0.6 to 1.4, simulating more or less forceful reps. Augmentations were randomly applied to the training windows on every retrieval.

C. Segmentation Model Design

As a first step toward achieving the goals of this work, a segmentation model was designed to detect the end of each rep. This model was designed to take in a window of 6-channel IMU data and output a binary classification of whether each input data point is the end of a rep. To train the model, each window was given 256 labels (one per data point) based on the end-rep-markers. The area (80ms) around any end-rep-markers occurring within a window was labelled with a positive label. The remaining data points of each window were labelled

with negative labels. An example of the labelling for a single window is shown in Figure 4.




Fig. 4. **Segmentation labels.** A single window of 1024 data points. A larger window size is chosen here for demonstration purposes. Shown from top to bottom are the acceleration data, gyroscope data, and segmentation labels. The segmentation labels are binary, with 1 indicating the end of a rep.

D. Classification Model Design

The next step in this work was to design a classification model to predict near-failure reps. This model was designed to classify a window of IMU data as either a near-failure rep ($RiR \leq 2$) or not. The model was trained on the same data as the segmentation model, but with a different set of labels. Each window was labelled with a binary label indicating whether the entire window was a near-failure rep or not. Windows with over 50% of their data points in a region of $RiR \leq 2$ (for

instance, the region corresponding to the last three reps in Figure 3) were given a positive label. The remaining windows were labelled with a negative label.

E. Model Architectures

1) *Segmentation Model*: The foundation of the segmentation model is a 1D convolutional residual neural network (ResNet) [10]. The model consists of several strided 1D convolutional layers, each followed by a batch normalization layer and a ReLU activation function. Residual connections are included around each convolutional layer. Convolutional layers with a kernel size of 1 are project residual connection channels when necessary. A final global average pooling and linear classification layer produce 256 output confidences, one for each data point in the input window. The final model architecture, found by neural architecture search (NAS) as explored in the next section, consists of 2,971,648 parameters and is shown in Figure 5 [11].




Fig. 5. **Segmentation model architecture.** Each block in the ResNet Base consists of a 1D Convolution layer with a kernel size of 3, followed by a batch normalization layer and a ReLU activation function. Residual connections are included around each block. Labels on each block represent the output channels and stride of each layer.

2) *Classification Model*: The base of the classification is the same ResNet base of the segmentation model. By loading the weights of the previous model, previously determined high dimensional representations of the IMU data can be used for further learning. This model also consists of an LSTM component




Fig. 6. **Classification model architecture.**

F. Neural Architecture Search

G. Model Training

H. Model Evaluation

I. Edge Deployment

describe steps taken to deploy to edge

III. RESULTS

Put any figures and tables here

A. Segmentation Model Performance

B. Classification Model Performance

C. Edge Deployment Performance

IV. DISCUSSION

V. AUTHOR CONTRIBUTIONS

A. Musa Azeem

Musa led the design and implementation of the models. He developed the initial architectures and refined them throughout the project. His responsibilities included model adaptation, training, and performance evaluation.

B. Grant King

Grant was responsible for collecting initial data from team members to debug and finalize the data collection tool. He also led the development of the deployment pipeline to the inference device and contributed to finalizing both the deployment process and the models during the later stages of the project.

C. Savannah Noblitt

Savannah designed and implemented the initial data pipeline and iteratively improved it over the course of the project. Her contributions included data cleaning, preparation, and augmentation to ensure high-quality inputs for training. She also contributed to finalizing both the data pipeline and the models during the later stages of the project.

REFERENCES

- [1] M. C. Refalo, E. R. Helms, Z. P. Robinson, D. L. Hamilton, and J. J. Fyfe, "Similar muscle hypertrophy following eight weeks of resistance training to momentary muscular failure or with repetitions-in-reserve in resistance-trained individuals," *Journal of sports sciences*, vol. 42, no. 1, pp. 85–101, 2024.
- [2] M. Izquierdo, J. Ibanez, J. J. González-Badillo, K. Häkkinen, N. A. Ratamess, W. J. Kraemer, D. N. French, J. Eslava, A. Altadill, X. Asiain *et al.*, "Differential effects of strength training leading to failure versus not to failure on hormonal responses, strength, and muscle power gains," *Journal of applied physiology*, 2006.
- [3] J. Kristensen and A. Franklyn-Miller, "Resistance training in musculoskeletal rehabilitation: a systematic review," *British journal of sports medicine*, vol. 46, no. 10, pp. 719–726, 2012.
- [4] J. A. Albert, A. Herdick, C. M. Brahms, U. Granacher, and B. Arnrich, "Using machine learning to predict perceived exertion during resistance training with wearable heart rate and movement sensors," in *2021 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*. IEEE, 2021, pp. 801–808.
- [5] Y. Jiang, V. Hernandez, G. Venture, D. Kulić, and B. K. Chen, "A data-driven approach to predict fatigue in exercise based on motion data from wearable sensors or force plate," *Sensors*, vol. 21, no. 4, p. 1499, 2021.
- [6] M. Elshafei and E. Shihab, "Towards detecting biceps muscle fatigue in gym activity using wearables," *Sensors*, vol. 21, no. 3, p. 759, 2021.

- [7] E. Velloso, A. Bulling, H. Gellersen, W. Ugulino, and H. Fuks, "Qualitative activity recognition of weight lifting exercises," in *Proceedings of the 4th Augmented Human International Conference*, 2013, pp. 116–123.
- [8] Y. Dang, Z. Liu, X. Yang, L. Ge, and S. Miao, "A fatigue assessment method based on attention mechanism and surface electromyography," *Internet of Things and Cyber-Physical Systems*, vol. 3, pp. 112–120, 2023.
- [9] K. Kings, "Best preacher curl alternatives," accessed April 30, 2025. [Online]. Available: <https://www.kettlebellkings.com/blogs/default-blog/best-preacher-curl-alternatives>
- [10] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 770–778.
- [11] B. Zoph and Q. V. Le, "Neural architecture search with reinforcement learning," *arXiv preprint arXiv:1611.01578*, 2016.